

▼ Homework 4

Instructions

- This homework focuses on understanding and applying CoCoOp for CLIP prompt tuning. It consists of **four questions** designed to assess both theoretical understanding and practical application.
- Please organize your answers and results for the questions below and submit this jupyter notebook as a **.pdf file**.
- **Deadline: 11/26 (Tue) 23:59**

➤ Preparation

- Run the code below before proceeding with the homework (Q1, Q2).
- If an error occurs, click 'Run Session Again' and then restart the runtime from the beginning.

[] ↩ 숨겨진 셀 1개

▼ Q1. Understanding and implementing CoCoOp

- We have learned how to define CoOp in Lab Session 4.
- The main difference between CoOp and CoCoOp is **meta network** to extract image tokens that is added to the text prompt.
- Based on the CoOp code given in Lab Session 4, fill-in-the-blank exercise to test your understanding of critical parts of the CoCoOp.

```

1 import torch.nn as nn
2
3 class CoCoOpPromptLearner(nn.Module):
4     def __init__(self, cfg, classnames, clip_model):
5         super().__init__()
6         n_cls = len(classnames)
7         n_ctx = cfg.TRAINER.COOCOOP.N_CTX
8         ctx_init = cfg.TRAINER.COOCOOP.CTX_INIT
9         dtype = clip_model.dtype
10        ctx_dim = clip_model.ln_final.weight.shape[0]
11        vis_dim = clip_model.visual.output_dim
12        clip_imsize = clip_model.visual.input_resolution
13        cfg_imsize = cfg.INPUT.SIZE[0]
14        assert cfg_imsize == clip_imsize, f"cfg_imsize ({cfg_imsize}) must equal to clip_imsize ({clip_imsize})"
15
16        if ctx_init:
17            # use given words to initialize context vectors
18            ctx_init = ctx_init.replace("_", " ")
19            n_ctx = len(ctx_init.split(" "))
20            prompt = clip.tokenize(ctx_init)
21            with torch.no_grad():
22                embedding = clip_model.token_embedding(prompt).type(dtype)
23            ctx_vectors = embedding[0, 1: 1 + n_ctx, :]
24            prompt_prefix = ctx_init
25        else:
26            # random initialization
27            ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)
28            nn.init.normal_(ctx_vectors, std=0.02)
29            prompt_prefix = " ".join(["X"] * n_ctx)
30
31        print(f'Initial context: "{prompt_prefix}"')
32        print(f"Number of context words (tokens): {n_ctx}")
33
34        self.ctx = nn.Parameter(ctx_vectors) # Wrap the initialized prompts above as parameters to make them trainable.
35
36        ### Tokenize ###
37        classnames = [name.replace("_", " ") for name in classnames] # 예) "Forest"
38        name_lens = [len(_tokenizer.encode(name)) for name in classnames]
39        prompts = [prompt_prefix + " " + name + "." for name in classnames] # 예) "A photo of Forest."
40
41        tokenized_prompts = torch.cat([clip.tokenize(p) for p in prompts]) # 예) [49406, 320, 1125, 539...]
42
43
44
45        #####
46        ##### Q1. Fill in the blank #####
47        ##### Define Meta Net #####
48        self.meta_net = nn.Sequential(OrderedDict([
49            ("linear1", nn.Linear(vis_dim, vis_dim//16)),
50            ("relu", nn.ReLU(inplace=True)),

```

```

51         ("linear2", nn.Linear(vis_dim // 16, ctx_dim))
52     ]))
53     #####
54     ## Hint: meta network is composed to linear layer, relu activation, and linear layer.
55
56
57
58     if cfg.TRAINER.COOCOOP.PREC == "fp16":
59         self.meta_net.half()
60
61     with torch.no_grad():
62         embedding = clip_model.token_embedding(tokenized_prompts).type(dtype)
63
64     # These token vectors will be saved when in save_model(),
65     # but they should be ignored in load_model() as we want to use
66     # those computed using the current class names
67     self.register_buffer("token_prefix", embedding[:, :1, :]) # SOS
68     self.register_buffer("token_suffix", embedding[:, 1 + n_ctx:, :]) # CLS, EOS
69     self.n_cls = n_cls
70     self.n_ctx = n_ctx
71     self.tokenized_prompts = tokenized_prompts # torch.Tensor
72     self.name_lens = name_lens
73
74     def construct_prompts(self, ctx, prefix, suffix, label=None):
75         # dim0 is either batch_size (during training) or n_cls (during testing)
76         # ctx: context tokens, with shape of (dim0, n_ctx, ctx_dim)
77         # prefix: the sos token, with shape of (n_cls, 1, ctx_dim)
78         # suffix: remaining tokens, with shape of (n_cls, *, ctx_dim)
79
80         if label is not None:
81             prefix = prefix[label]
82             suffix = suffix[label]
83
84         prompts = torch.cat(
85             [
86                 prefix, # (dim0, 1, dim)
87                 ctx, # (dim0, n_ctx, dim)
88                 suffix, # (dim0, *, dim)
89             ],
90             dim=1,
91         )
92
93         return prompts
94
95     def forward(self, im_features):
96         prefix = self.token_prefix
97         suffix = self.token_suffix
98         ctx = self.ctx # (n_ctx, ctx_dim)
99
100
101
102     #####
103     ##### Q2.3. Fill in the blank #####
104     bias=self.meta_net(im_features) # (batch, ctx_dim)
105     bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)
106     ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
107     ctx_shifted=ctx+bias # (batch, n_ctx, ctx_dim)
108     #####
109     #####
110
111
112     # Use instance-conditioned context tokens for all classes
113     prompts = []
114     for ctx_shifted_i in ctx_shifted:
115         ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
116         pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
117         prompts.append(pts_i)
118     prompts = torch.stack(prompts)
119
120
121     return prompts

```

```

1 class CoCoOpCustomCLIP(nn.Module):
2     def __init__(self, cfg, classnames, clip_model):
3         super().__init__()
4         self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
5         self.tokenized_prompts = self.prompt_learner.tokenized_prompts
6         self.image_encoder = clip_model.visual
7         self.text_encoder = TextEncoder(clip_model)
8         self.logit_scale = clip_model.logit_scale
9         self.dtype = clip_model.dtype

```

```

11 def forward(self, image, label=None):
12     tokenized_prompts = self.tokenized_prompts
13     logit_scale = self.logit_scale.exp()
14
15     image_features = self.image_encoder(image.type(self.dtype))
16     image_features = image_features / image_features.norm(dim=-1, keepdim=True)
17
18
19     #####
20     ##### Q4. Fill in the blank #####
21     prompts=self.prompt_learner(image_features)
22     #####
23     #####
24
25
26     logits = []
27     for pts_i, imf_i in zip(prompts, image_features):
28         text_features = self.text_encoder(pts_i, tokenized_prompts)
29         text_features = text_features / text_features.norm(dim=-1, keepdim=True)
30         l_i = logit_scale * imf_i @ text_features.t()
31         logits.append(l_i)
32     logits = torch.stack(logits)
33
34     if self.prompt_learner.training:
35         return F.cross_entropy(logits, label)
36
37     return logits

```

✓ Q2. Training CoCoOp

In this task, you will train CoCoOp on the EuroSAT dataset. If your implementation of CoCoOp in Question 1 is correct, the following code should execute without errors. Please submit the execution file so we can evaluate whether your code runs without any issues.

```

1 # Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.
2 args.trainer = "CoCoOp"
3 args.train_batch_size = 4
4 args.epoch = 100
5 args.output_dir = "outputs/cocoop"
6
7 args.subsample_classes = "base"
8 args.eval_only = False
9 cocoop_base_acc = main(args)

```

```

➡ Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Creating a 16-shot dataset
Creating a 4-shot dataset
Saving preprocessed few-shot data to /content/ProMetaR/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE BASE CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 8 worker processes in total.
warnings.warn(

Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       4,200

Loading CLIP (backbone: ViT-B/16)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.linear2.weight', 'prompt_learner.meta_net.linear2.bias', 'prompt_learner.meta_net.linear1}
Loading evaluator: Classification
No checkpoint found, train from scratch
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_
warnings.warn(
epoch [1/100] batch [20/20] time 0.100 (0.333) data 0.000 (0.047) loss 0.2744 (1.1881) lr 2.5000e-03 eta 0:10:58
epoch [2/100] batch [20/20] time 0.091 (0.123) data 0.000 (0.017) loss 0.8384 (0.8970) lr 2.4994e-03 eta 0:04:00
epoch [3/100] batch [20/20] time 0.092 (0.122) data 0.000 (0.020) loss 0.6382 (0.7859) lr 2.4975e-03 eta 0:03:55
epoch [4/100] batch [20/20] time 0.092 (0.120) data 0.000 (0.023) loss 0.5044 (0.7151) lr 2.4945e-03 eta 0:03:50

```

```

epoch [5/100] batch [20/20] time 0.136 (0.188) data 0.000 (0.034) loss 0.5703 (0.6317) lr 2.4901e-03 eta 0:05:56
epoch [6/100] batch [20/20] time 0.090 (0.124) data 0.000 (0.025) loss 0.6060 (0.6009) lr 2.4846e-03 eta 0:03:52
epoch [7/100] batch [20/20] time 0.090 (0.141) data 0.000 (0.016) loss 0.3853 (0.6638) lr 2.4779e-03 eta 0:04:21
epoch [8/100] batch [20/20] time 0.092 (0.123) data 0.000 (0.017) loss 1.4082 (0.6633) lr 2.4699e-03 eta 0:03:46
epoch [9/100] batch [20/20] time 0.120 (0.137) data 0.000 (0.019) loss 0.1780 (0.4582) lr 2.4607e-03 eta 0:04:09
epoch [10/100] batch [20/20] time 0.136 (0.209) data 0.000 (0.033) loss 1.2285 (0.5051) lr 2.4504e-03 eta 0:06:16
epoch [11/100] batch [20/20] time 0.094 (0.130) data 0.000 (0.017) loss 0.2539 (0.5013) lr 2.4388e-03 eta 0:03:50
epoch [12/100] batch [20/20] time 0.091 (0.124) data 0.000 (0.021) loss 1.1484 (0.4657) lr 2.4261e-03 eta 0:03:38
epoch [13/100] batch [20/20] time 0.100 (0.122) data 0.000 (0.015) loss 0.8467 (0.5009) lr 2.4122e-03 eta 0:03:32
epoch [14/100] batch [20/20] time 0.129 (0.137) data 0.000 (0.019) loss 0.5547 (0.4495) lr 2.3972e-03 eta 0:03:56
epoch [15/100] batch [20/20] time 0.134 (0.189) data 0.000 (0.032) loss 1.0430 (0.5549) lr 2.3810e-03 eta 0:05:21
epoch [16/100] batch [20/20] time 0.092 (0.123) data 0.000 (0.018) loss 1.3906 (0.4799) lr 2.3638e-03 eta 0:03:26
epoch [17/100] batch [20/20] time 0.090 (0.122) data 0.000 (0.019) loss 0.0238 (0.3497) lr 2.3454e-03 eta 0:03:22
epoch [18/100] batch [20/20] time 0.095 (0.121) data 0.000 (0.019) loss 0.1337 (0.2804) lr 2.3259e-03 eta 0:03:18
epoch [19/100] batch [20/20] time 0.134 (0.133) data 0.000 (0.017) loss 1.0420 (0.3864) lr 2.3054e-03 eta 0:03:35
epoch [20/100] batch [20/20] time 0.139 (0.184) data 0.000 (0.025) loss 0.3484 (0.4984) lr 2.2839e-03 eta 0:04:54
epoch [21/100] batch [20/20] time 0.094 (0.122) data 0.000 (0.016) loss 0.8184 (0.3424) lr 2.2612e-03 eta 0:03:20

1 # Accuracy on the New Classes.
2 args.model_dir = "outputs/cocoop"
3 args.output_dir = "outputs/cocoop/new_classes"
4 args.subsample_classes = "new"
5 args.load_epoch = 100
6 args.eval_only = True
7 cocoop_novel_acc = main(args)

Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])

-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       3,900
-----

Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 8 worker processes in total
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr
  warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value)
  checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.linear2.bias', 'prompt_learner.meta_net.linear2.weight', 'prompt_learner.meta_net.linear1.we
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 100)
Evaluate on the *test* set
100%|██████████| 39/39 [01:04<00:00, 1.66s/it]=> result
* total: 3,900
* correct: 1,687
* accuracy: 43.3%
* error: 56.7%
* macro_f1: 39.0%

```

✓ Q3. Analyzing the results of CoCoOp

Compare the results of CoCoOp with those of CoOp that we trained in Lab Session 4. Discuss possible reasons for the performance differences observed between CoCoOp and CoOp.

CoOp and CoCoOp both investigate the most efficient input phrases instead of fixing format like 'a photo of {object}'.

However, while CoOp focuses solely on phrases and classes, CoCoOp utilizes information from the image file. It prevents overfitting problem in CoOp, and this is implemented at the meta network; the homework part of Q1.

The differences can result in following consequences:

- CoOp gets higher accuracy when faced on base classes.
- CoCoOp gets higher accuracy when faced on unseen classes than CoOp.
- In overall situation, CoCoOp shows better performance.

Comparing the above with the result of CoOp in the Lab session 4, it performs worse in all situations. ($90.8\% < 92.3\%$, $43.3\% < 51.5\%$)

Nevertheless, it does not mean that it is a counterexample against the claim. In this project, it sets learnable context as 4 words, especially regarding 'a photo of a' as a default state. On the other hand, CoOp implemented in Lab session 4 suggests learnable context as 16 words and it starts from none-state (blank).

Meanwhile, the speed of CoOp is faster than CoCoOp. This can be seen in runtime comparison.